

Text Representation for Classification


LÊ ANH CƯỜNG

Faculty of Information Technology


Ton Duc Thang University, Ho Chi Minh city, Vietnam

2020


Natural Language Processing: a General Scheme



Natural Language Processing: a General Scheme



Natural Language Processing: a General Scheme



Tasks with inputs/outputs

Task	Input	output
Document classification	A textual document	Classes/labels
Parsing	A sentence	Parsed tree
Machine translation	A source sentence	A target sentence
Named entity recognition	A sentence	NE tags
...

How represent the **Input** for NLP tasks?

- A sentence: a sequence of Words
- A document: a sequence of sentences
- A word: a sequence of characters

- The **word** is the smallest unit of language that has a meaning

Work on Documents?

- Document representation
- Document generation
- Document classification
- Document summarization, multi-document summarization

Types of Word Representation

1. Word as unique, indivisible unit
2. Word as Distributed Semantic Representation

Word as unique, indivisible unit

- Each word is represented as *this word*
- Each word is represented by an unique **key** or an **index** in the vocabulary dictionary.
- Each word is represented as **one-hot vector**

Rome = [1, 0, 0, 0, 0, 0, ..., 0]

Paris = [0, 1, 0, 0, 0, 0, ..., 0]

Italy = [0, 0, 1, 0, 0, 0, ..., 0]

France = [0, 0, 0, 1, 0, 0, ..., 0]

word V = [0, 0, 0, 0, 0, 0, ..., 0]

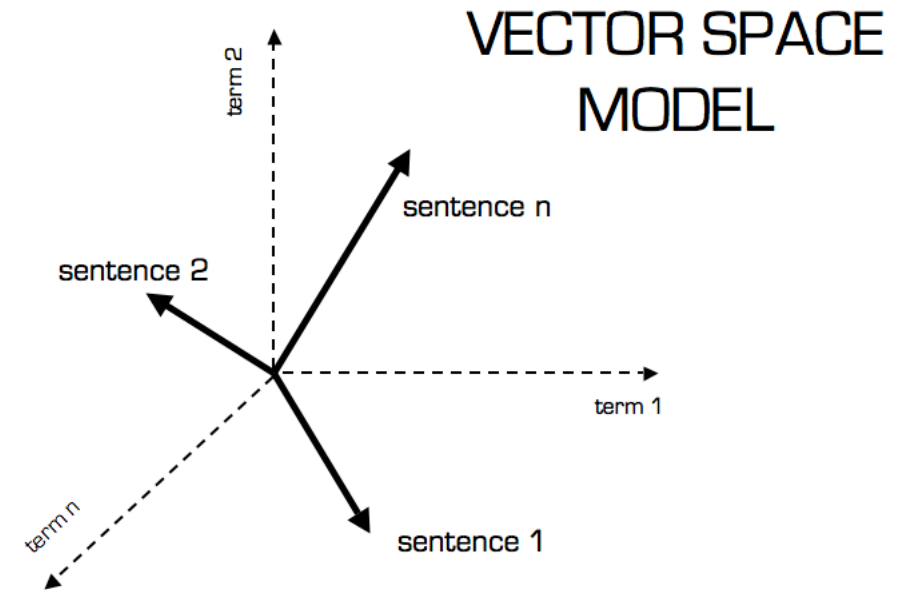

Text Representation

		Terms			
Documents		data	result	statistics	analysis
	Document1	0	1	0	1
	Document2	1	0	1	0
	Document3	0	0	1	0
	Document4	1	1	0	0

	it	is	puppy	cat	pen	a	this
it is a puppy	1	1	1	0	0	1	0
it is a kitten	1	1	0	0	0	1	0
it is a cat	1	1	0	1	0	1	0
that is a dog and this is a pen	0	2	0	0	1	2	1
it is a matrix	1	1	0	0	0	1	0

- Represent each document as a feature vector of Vector Space Model (Term Vector Model)
- **Term**: word, n-gram, ...

Vector Space Model and Bag-Of-Words for Text Representation



TF-IDF weights

$$w_{i,j} = tf_{i,j} \times \log \left(\frac{N}{df_i} \right)$$

tf_{ij} = number of occurrences of i in j
 df_i = number of documents containing i
 N = total number of documents

Variants of TF weight

weighting scheme	TF weight
binary	{0,1}
raw frequency	$f_{t,d}$
log normalization	$\log(1 + f_{t,d})$
double normalization 0.5	$0.5 + 0.5 \frac{f_{t,d}}{\max f_{t,d}}$
double normalization K	$K + (1 - K) \frac{f_{t,d}}{\max f_{t,d}}$

TF-IDF example

Document-term matrix (no need to normalize, every word occurs just once)

	angeles	los	new	post	times	york
d1	0	0	1	0	1	1
d2	0	0	1	1	0	1
d3	1	1	0	0	1	0

tf-idf

	angeles	los	new	post	times	york
d1	0	0	0.584	0	0.584	0.584
d2	0	0	0.584	1.584	0	0.584
d3	1.584	1.584	0	0	0.584	0

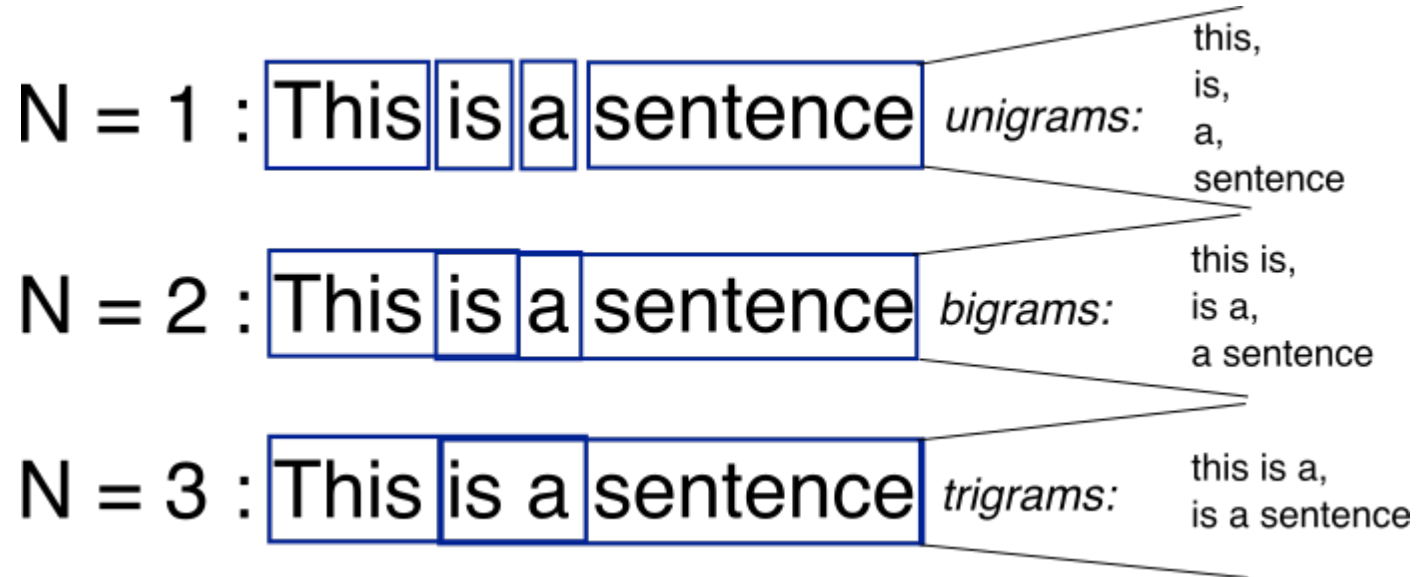
TF-IDF Variants

Term frequency		Document frequency		Normalization	
n (natural)	$tf_{t,d}$	n (no)	1	n (none)	1
l (logarithm)	$1 + \log(tf_{t,d})$	t (idf)	$\log \frac{N}{df_t}$	c (cosine)	$\frac{1}{\sqrt{w_1^2 + w_2^2 + \dots + w_M^2}}$
a (augmented)	$0.5 + \frac{0.5 \times tf_{t,d}}{\max_t(tf_{t,d})}$	p (prob idf)	$\max\{0, \log \frac{N - df_t}{df_t}\}$	u (pivoted unique)	$1/u$ (Section 6.4.4)
b (boolean)	$\begin{cases} 1 & \text{if } tf_{t,d} > 0 \\ 0 & \text{otherwise} \end{cases}$			b (byte size)	$1/CharLength^\alpha, \alpha < 1$
L (log ave)	$\frac{1 + \log(tf_{t,d})}{1 + \log(\text{ave}_{t \in d}(tf_{t,d}))}$				

► **Figure 6.7** SMART notation for tf-idf variants. Here *CharLength* is the number of characters in the document.

Terms ?

- N-gram



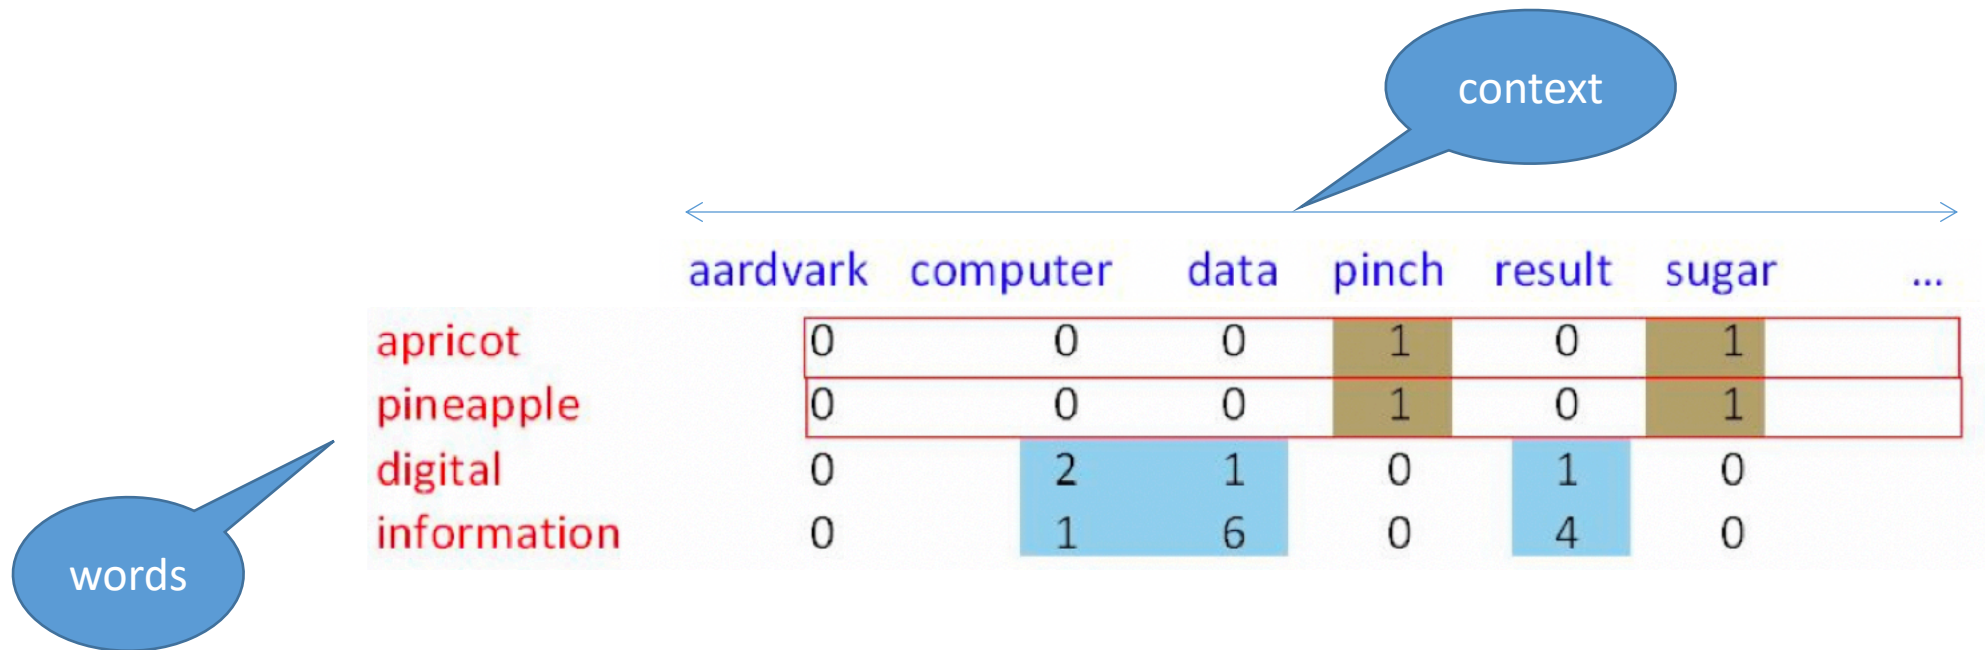
- Other kinds: POS tags, Syntactic sub-structures, ...

Applications

- Text Similarity
- Text Classification
- Word Similarity
- Information Retrieval

Term-context matrix for word similarity

- Two words are similar in meaning if their context vector are similar



The diagram illustrates a term-context matrix. A blue callout labeled 'words' points to the first column of the matrix, which lists the words: apricot, pineapple, digital, and information. Another blue callout labeled 'context' points to the top row of the matrix, which lists the context words: aardvark, computer, data, pinch, result, sugar, and The matrix itself is a table where rows represent words and columns represent context words. The cells contain numerical values representing the relationship between the word and the context. Some cells are highlighted in brown (1), blue (2, 1, 4, 6), or white (0).

	aardvark	computer	data	pinch	result	sugar	...
apricot	0	0	0	1	0	1	
pineapple	0	0	0	1	0	1	
digital	0	2	1	0	1	0	
information	0	1	6	0	4	0	

Cosine Similarity for Information Retrieval, as well as for Text Similarity


Documents and queries are represented as vectors.

$$d_j = (w_{1,j}, w_{2,j}, \dots, w_{t,j})$$

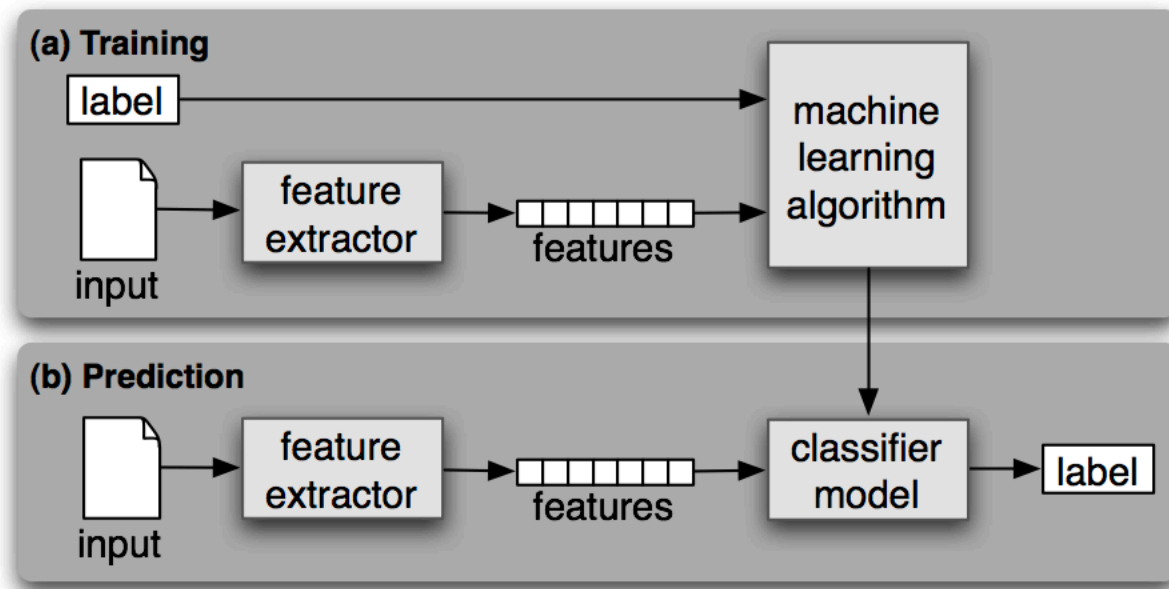
$$q = (w_{1,q}, w_{2,q}, \dots, w_{n,q})$$

Using the cosine the similarity between document d_j and query q can be calculated as:


$$\cos(d_j, q) = \frac{\mathbf{d}_j \cdot \mathbf{q}}{\|\mathbf{d}_j\| \|\mathbf{q}\|} = \frac{\sum_{i=1}^N w_{i,j} w_{i,q}}{\sqrt{\sum_{i=1}^N w_{i,j}^2} \sqrt{\sum_{i=1}^N w_{i,q}^2}}$$



Machine Learning for Text Classification



Machine Learning for Text Classification




Projects

1. Document classification
2. Word Similarity

- Requirement:
 - Using “vector space model”
 - Using different representations of document and word context:
 - Binary matrix
 - Word Frequency
 - tf.idf
 - N-gram

Text preprocessing

- Token segmentation/split
- Sentence segmentation
- Get n-gram of token sequences
- Vocabulary of tokens
- Texts to sequences
- Sequences to texts
- Convert sentence to vector
 - Binary mode
 - Counting mode
 - Tfidf mode



Text Preprocessing: using nltk

Input: document d

- Sentence segmentation: s_1, s_2, \dots, s_n
- For each sentence s_i :
- Token segmentation: w_1, w_2, \dots, w_k
- Vocabulary dictionary
- Convert word to number
- Convert sequence (i.e.) to vector

```
>>> from nltk import tokenize
>>> p = "Good morning Dr. Adams. The patient is waiting for you in room number 123."

>>> tokenize.sent_tokenize(p)
['Good morning Dr. Adams.', 'The patient is waiting for you in room number 123.']
```

```
>>> import nltk
>>> sentence = """At eight o'clock on Thursday morning
... Arthur didn't feel very good."""
>>> tokens = nltk.word_tokenize(sentence)
>>> tokens
['At', 'eight', 'o'clock', 'on', 'Thursday', 'morning',
'Arthur', 'did', 'n't', 'feel', 'very', 'good', '.']
```

```
from nltk.tokenize import word_tokenize
from gensim.corpora.dictionary import Dictionary

sometext = "hello how are you doing?"

tokens = word_tokenize(sometext)
my_vocab = Dictionary([tokens])

print(my_vocab.token2id['hello'])
```



```
from nltk.util import ngrams
text = "Hi How are you? i am fine and you"
n = int(input("ngram value = "))
n_grams = ngrams(text.split(), n)
for grams in n_grams :
    print(grams)
```

Text preprocessing: using tensorflow.keras.preprocessing.text

https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/text/Tokenizer

▼ text

- Overview
- hashing_trick
- one_hot
- text_to_word_sequence
- Tokenizer**
- tokenizer_from_json

▼ text

- Overview
- hashing_trick
- one_hot
- text_to_word_sequence**
- Tokenizer
- tokenizer_from_json

tf.keras.preprocessing.text.Tokenizer

<https://keras.io/api/preprocessing/text/>

Tokenizer class

```
tf.keras.preprocessing.text.Tokenizer(  
    num_words=None,  
    filters='!"#$%&()*+,-./:;<=>?@[\\]^_`{|}~\t\n',  
    lower=True,  
    split=" ",  
    char_level=False,  
    oov_token=None,  
    document_count=0,  
    **kwargs  
)
```

tf.keras.preprocessing.text.Tokenizer

Tokenizer class

```
tf.keras.preprocessing.text.Tokenizer(  
    num_words=None,  
    filters='!"#$%&()*+,-./:;<=>?@[\\]^_`{|}~\t\n',  
    lower=True,  
    split=" ",  
    char_level=False,  
    oov_token=None,  
    document_count=0,  
    **kwargs  
)
```

- **num_words**: the maximum number of words to keep, based on word frequency. Only the most common `num_words-1` words will be kept.
- **filters**: a string where each element is a character that will be filtered from the texts. The default is all punctuation, plus tabs and line breaks, minus the ' ' character.
- **lower**: boolean. Whether to convert the texts to lowercase.
- **split**: str. Separator for word splitting.
- **char_level**: if True, every character will be treated as a token.
- **oov_token**: if given, it will be added to word_index and used to replace out-of-vocabulary words during text_to_sequence calls

tf.keras.preprocessing.text.Tokenizer

```
from keras.preprocessing.text import Tokenizer
```

Using TensorFlow backend.

```
t = Tokenizer()
text = 'The earth is an awesome place live'
tokens = text.split(' ')
t.fit_on_texts(tokens)
print("word_index : ",t.word_index)
```

```
word_index : {'the': 1, 'earth': 2, 'is': 3, 'an': 4, 'awesome': 5, 'place': 6, 'live': 7}
```

```
test_text = ['The', 'earth', 'live']
sequences = t.texts_to_sequences(test_text)
print('sequences : ',sequences,'\n')
```

```
sequences : [[1], [2], [7]]
```

tf.keras.preprocessing.text.Tokenizer

```
from keras.preprocessing.text import Tokenizer
```

Using TensorFlow backend.

```
t = Tokenizer()  
text = ['The earth is an awesome place live']  
t.fit_on_texts(text)  
print("word_index : ",t.word_index)
```

```
word_index : {'the': 1, 'earth': 2, 'is': 3, 'an': 4, 'awesome': 5, 'place': 6, 'live': 7}
```

```
test_text = ['The', 'earth', 'live']  
sequences = t.texts_to_sequences(test_text)  
print('sequences : ',sequences,'\n')
```

```
sequences : [[1], [2], [7]]
```

tf.keras.preprocessing.text.Tokenizer

```
docs = ['Well done!',  
        'Good work',  
        'Great effort',  
        'nice work',  
        'Excellent!']  
  
# create the tokenizer  
t = Tokenizer()  
# fit the tokenizer on the documents  
t.fit_on_texts(docs)  
t.word_index.items()
```

```
dict_items([('work', 1), ('well', 2), ('done', 3), ('good', 4), ('great', 5), ('effort', 6),  
           ('nice', 7), ('excellent', 8)])
```

```
test_text = ['Good nice work']  
sequences = t.texts_to_sequences(test_text)  
print('sequences : ', sequences, '\n')
```

```
sequences :  [[4, 7, 1]]
```

tf.keras.preprocessing.text.Tokenizer sequence to texts

```
test_text = ['Good nice work']  
sequences = t.texts_to_sequences(test_text)  
print('sequences : ', sequences, '\n')
```

```
sequences :  [[4, 7, 1]]
```

```
text = t.sequences_to_texts(sequences)  
print(text)
```

```
['good nice work']
```

```
text = t.sequences_to_texts([[7]])  
print(text)
```

```
['nice']
```


tf.keras.preprocessing.text. **text_to_word_sequence**

```
tf.keras.preprocessing.text.text_to_word_sequence(  
    text, filters='!"#$%&()*+,-./:;<=>?@[\\]^_`{|}~\t\n', lower=True, split=' '  
)
```

tf.keras.preprocessing.text. text_to_word_sequence

```
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.text import text_to_word_sequence
max_words = 10000

text = 'Decreased glucose-6-phosphate dehydrogenase activity along with oxidative stress affects visual contrast sensitivity in alcoholics'
text = text_to_word_sequence(text)
print(text)
tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(text)
sequences = tokenizer.texts_to_sequences(text)
print(sequences)
```

```
[['decreased', 'glucose', '6', 'phosphate', 'dehydrogenase', 'activity', 'along', 'with', 'oxidative', 'stress', 'affects', 'visual', 'contrast', 'sensitivity', 'in', 'alcoholics']]
[[1], [2], [3], [4], [5], [6], [7], [8], [9], [10], [11], [12], [13], [14], [15], [16]]
```

Convert to vector

`fit_on_texts`

```
fit_on_texts(  
    texts  
)
```

Updates internal vocabulary based on a list of texts.

In the case where texts contains lists, we assume each entry of the lists to be a token.

Required before using `texts_to_sequences` or `texts_to_matrix`.

`texts_to_matrix`

```
texts_to_matrix(  
    texts, mode='binary'  
)
```

Convert a list of texts to a Numpy matrix.

Arguments

```
texts: list of strings.  
mode: one of "binary", "count", "tfidf", "freq".
```

tf.keras.preprocessing.text.Tokenizer

```
docs = ['Well done!',  
        'Good work',  
        'Great effort',  
        'nice work',  
        'Excellent!']  
  
# create the tokenizer  
t = Tokenizer()  
  
# fit the tokenizer on the documents  
t.fit_on_texts(docs)  
print(t)  
encoded_docs = t.texts_to_matrix(docs, mode='count')  
print(encoded_docs)  
print(t.word_index.items())
```

```
[[0. 0. 1. 1. 0. 0. 0. 0. 0.]  
 [0. 1. 0. 0. 1. 0. 0. 0. 0.]  
 [0. 0. 0. 0. 0. 1. 1. 0. 0.]  
 [0. 1. 0. 0. 0. 0. 0. 1. 0.]  
 [0. 0. 0. 0. 0. 0. 0. 0. 1.]]
```

tf.keras.preprocessing.text.Tokenizer

tf.idf

```
docs = ['Well done!',  
        'Good work',  
        'Great effort',  
        'nice work',  
        'Excellent!']  
  
# create the tokenizer  
t = Tokenizer()  
# fit the tokenizer on the documents  
t.fit_on_texts(docs)  
encoded_docs = t.texts_to_matrix(docs, mode='tfidf')  
print(encoded_docs)  
print(t.word_index.items())
```

```
[[0.          0.          1.25276297 1.25276297 0.          0.  
  0.          0.          0.          ]  
 [0.          0.98082925 0.          0.          1.25276297 0.  
  0.          0.          0.          ]  
 [0.          0.          0.          0.          0.          1.25276297  
  1.25276297 0.          0.          ]  
 [0.          0.98082925 0.          0.          0.          0.  
  0.          1.25276297 0.          ]  
 [0.          0.          0.          0.          0.          0.  
  0.          0.          1.25276297]]
```

```
dict_items([('work', 1), ('well', 2), ('done', 3), ('good', 4), ('great', 5), ('effort', 6),  
           ('nice', 7), ('excellent', 8)])
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
vectorizer = TfidfVectorizer(max_features=300)
vectorizer = vectorizer.fit(X_train)

df_train = vectorizer.transform(X_train)

print(df_train.shape)
print(df_train)
```

```
(900, 300)
(0, 133)      1.0
(1, 89)       0.7403480856124964
(1, 65)       0.6722237069085794
(2, 246)      0.30211756738541284
(2, 238)      0.6850871121270957
(2, 100)      0.4174660000455605
```

```
from sklearn.feature_extraction.text import TfidfTransformer  
from sklearn.feature_extraction.text import CountVectorizer
```

<https://kavita-ganesan.com/tfidftransformer-tfidfvectorizer-usage-differences/#.Xsij-RMzbEY>